configuring one more as worker node..

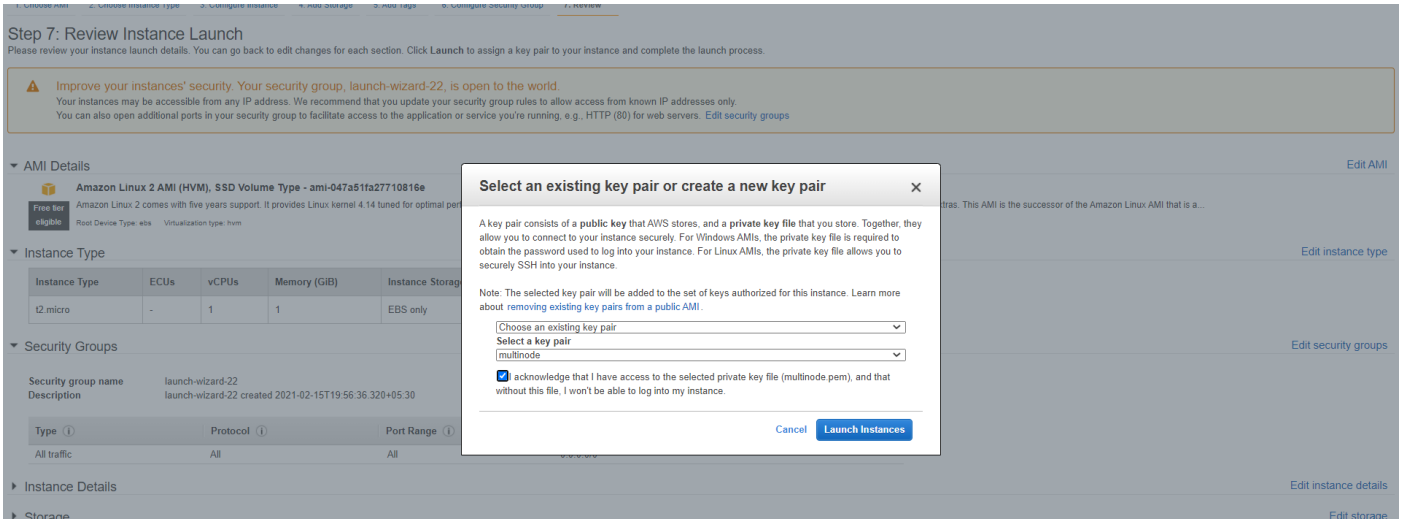

launching one more….

⇨ connect to putty

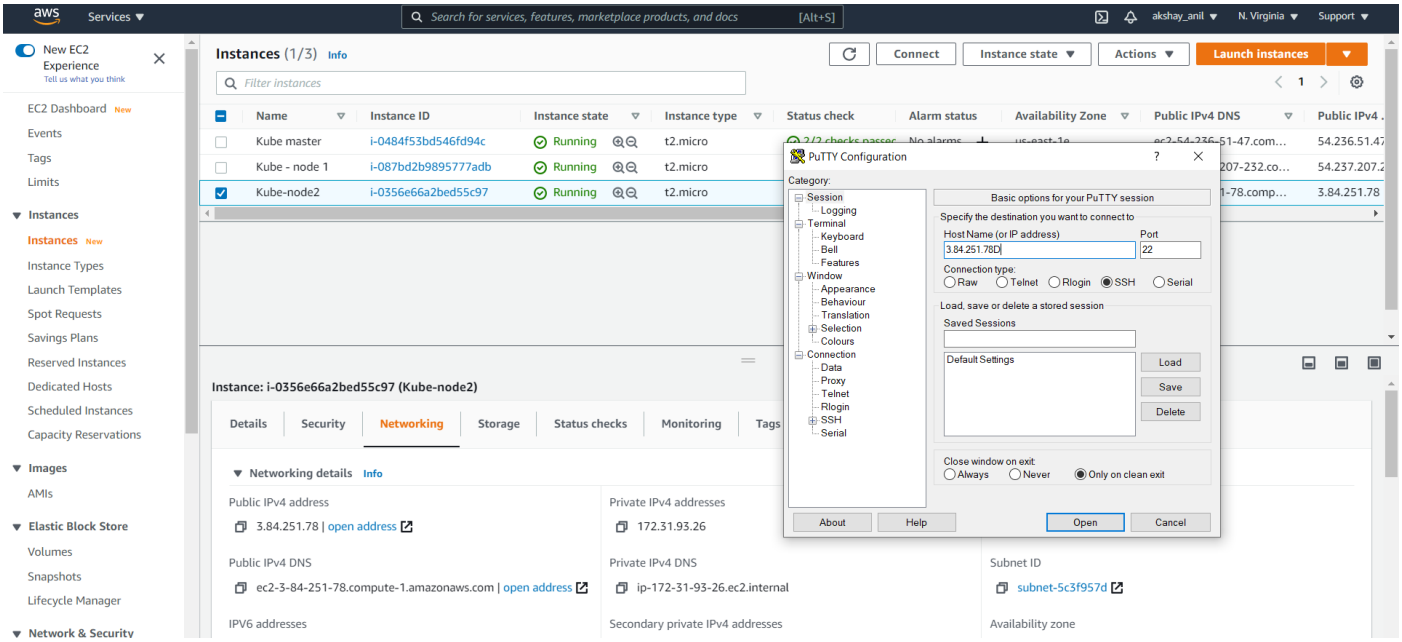

⇨ login into root

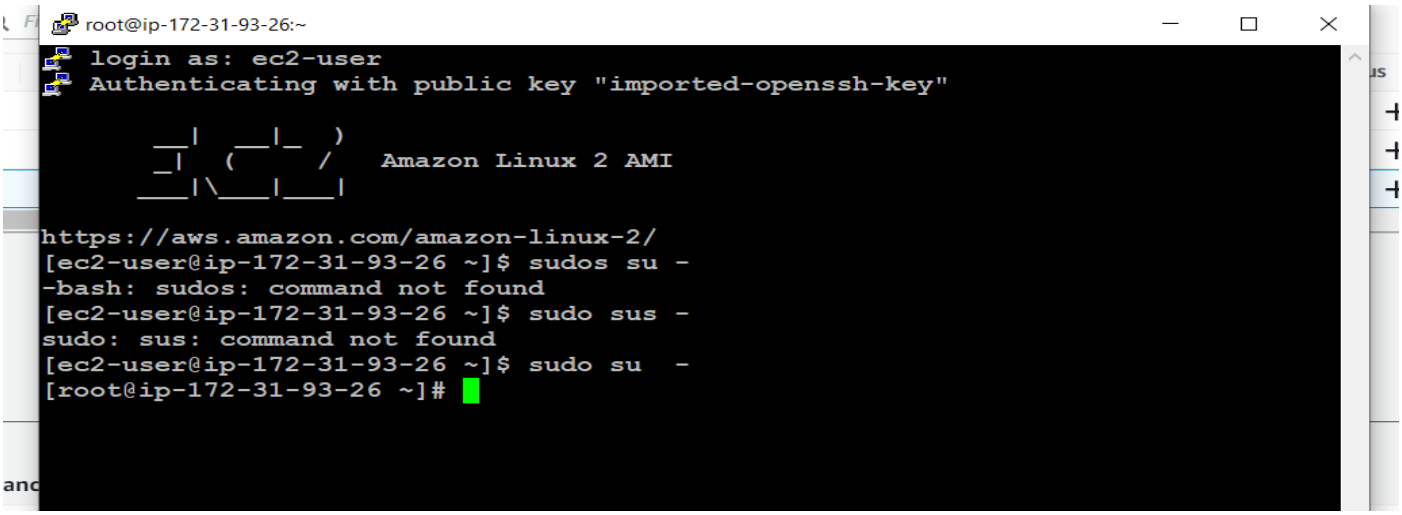

installation of docker

1. yum install docker -y
2. systemctl enable docker --now
3. docker info

```
[root@ip-172-31-93-26 ~]# yum install docker -y
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
Resolving Dependencies
```

```
[root@ip-172-31-93-26 ~]# systemctl enable docker --now
Created symlink from /etc/systemd/system/multi-user.target.wants/docker.service
to /usr/lib/systemd/system/docker.service.
[root@ip-172-31-93-26 ~]#
```

```
[root@ip-172-31-93-26 ~]# docker info
Client:
 Debug Mode: false

Server:
 Containers: 0
  Running: 0
  Paused: 0
  Stopped: 0
 Images: 0
 Server Version: 19.03.13-ce
 Storage Driver: overlay2
  Backing Filesystem: xfs
  Supports d_type: true
  Native Overlay Diff: true
 Logging Driver: json-file
 Cgroup Driver: cgroupfs
 Plugins:
  Volume: local
```

installation of kubeadm

1. yum install kubeadm - fails
2. vi /etc/yum.repos.d/kubernetes.repo - need to setup the repo
3. yum repolist
4. yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
5. systemctl status kubelet
6. systemctl enable kubelet –now

```
[root@ip-172-31-93-26 ~]# vi /etc/yum.repos.d/kubernetes.repo
[root@ip-172-31-93-26 ~]# cat  /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-\$basearch
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
exclude=kubelet kubeadm kubectl
[root@ip-172-31-93-26 ~]#
```

```
[root@ip-172-31-93-26 ~]# yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
Resolving Dependencies
--> Running transaction check
---> Package kubeadm.x86_64 0:1.20.2-0 will be installed
--> Processing Dependency: kubernetes-cni >= 0.8.6 for package: kubeadm-1.20.2-0.x86_64
--> Processing Dependency: cri-tools >= 1.13.0 for package: kubeadm-1.20.2-0.x86_64
---> Package kubectl.x86_64 0:1.20.2-0 will be installed
---> Package kubelet.x86_64 0:1.20.2-0 will be installed
--> Processing Dependency: socat for package: kubelet-1.20.2-0.x86_64
--> Processing Dependency: ebtables for package: kubelet-1.20.2-0.x86_64
--> Processing Dependency: conntrack for package: kubelet-1.20.2-0.x86_64
--> Running transaction check
[root@ip-172-31-93-26 ~]# systemctl enable kubelet --now
Created symlink from /etc/systemd/system/multi-user.target.wants/kubelet.service to /usr/lib/systemd/system/kubelet.service.
[root@ip-172-31-93-26 ~]#
```

```
[root@ip-172-31-93-26 ~]# systemctl status kubelet
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/usr/lib/systemd/system/kubelet.service; enabled; vendor preset: disabled)
  Drop-In: /usr/lib/systemd/system/kubelet.service.d
           └─10-kubeadm.conf
   Active: activating (auto-restart) (Result: exit-code) since Mon 2021-02-15 14:46:46 UTC; 2s ago
     Docs: https://kubernetes.io/docs/
  Process: 4250 ExecStart=/usr/bin/kubelet $KUBELET_KUBECONFIG_ARGS $KUBELET_CONFIG_ARGS $KUBELET_KUBEADM_ARGS $KUBELET_EXTRA_ARGS (code=exited, status=255)
 Main PID: 4250 (code=exited, status=255)

Feb 15 14:46:46 ip-172-31-93-26.ec2.internal kubelet[4250]: created by k8s.io/kubernetes/vendor/go.opencensus.io/stats/view.init.0
```

..

now, we have to join this to master

Kubeadm join.

```
Hint: Some lines were ellipsized, use -l to show in full.
[root@ip-172-31-93-26 ~]# kubeadm join
discovery: Invalid value: "": bootstrapToken or file must be set
To see the stack trace of this error execute with --v=5 or higher
[root@ip-172-31-93-26 ~]#
```

right now, it fails, it requires token from master,

since we know, preflight check error will come,

1st error: [WARNING IsDockerSystemdCheck]: we have to use system instead of cgroups.

vi /etc/docker/daemon.json

```
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
```

cat /etc/docker/daemon.json

systemctl restart docker

docker info | grep Driver

```
[root@ip-172-31-93-26 ~]# vi /etc/docker/daemon.json
[root@ip-172-31-93-26 ~]# cat /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
[root@ip-172-31-93-26 ~]# systemctl restart docker
[root@ip-172-31-93-26 ~]# docker info | grep driver
[root@ip-172-31-93-26 ~]# docker info | grep Driver
 Storage Driver: overlay2
 Logging Driver: json-file
 Cgroup Driver: systemd
[root@ip-172-31-93-26 ~]#
```

2nd error: [WARNING FileExisting-tc]: tc not found in system path

yum install -y iproute-tc

```
[root@ip-172-31-93-26 ~]# yum install -y iproute-tc
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
amzn2-core                                                                | 3.7 kB  00:00:00
Resolving Dependencies
--> Running transaction check
---> Package iproute-tc.x86_64 0:4.15.0-1.amzn2.0.4 will be installed
--> Finished Dependency Resolution

Dependencies Resolved
```

3rd error: FileContent—proc-sys-net-bridge-bridge-nf-call-iptables

vim /etc/sysctl.d/k8s.conf

net.bridge.bridge-nf-call-ip6tables = 1

net.bridge.bridge-nf-call-iptables = 1

cat /etc/sysctl.d/k8s.conf

sysctl –system

sysctl -a | grep bridge-bridge-nf-call

```
[root@ip-172-31-93-26 ~]# vim /etc/sysctl.d/k8s.conf
[root@ip-172-31-93-26 ~]# cat /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-ip6tables = 1
net.bridge.bridge-nf-call-iptables = 1
```

```
[root@ip-172-31-93-26 ~]# sysctl --system
* Applying /etc/sysctl.d/00-defaults.conf ...
kernel.printk = 8 4 1 7
kernel.panic = 30
net.ipv4.neigh.default.gc_thresh1 = 0
net.ipv6.neigh.default.gc_thresh1 = 0
net.ipv4.neigh.default.gc_thresh2 = 15360
net.ipv6.neigh.default.gc_thresh2 = 15360
net.ipv4.neigh.default.gc_thresh3 = 16384
* Applying /etc/sysctl.conf ...
[root@ip-172-31-93-26 ~]# sysctl -a | grep bridge-bridge-nf-call
sysctl: reading key "net.ipv6.conf.all.stable_secret"
sysctl: reading key "net.ipv6.conf.default.stable_secret"
sysctl: reading key "net.ipv6.conf.docker0.stable_secret"
sysctl: reading key "net.ipv6.conf.eth0.stable_secret"
sysctl: reading key "net.ipv6.conf.lo.stable_secret"
[root@ip-172-31-93-26 ~]#
```

now, we can join this worker to the master

go to master create the token

```
Last login: Mon Feb 15 14:08:30 2021 from ec2-18-206-107-24.compute-1.amazonaws.com

       __|  __|_  )
       _|  (     /   Amazon Linux 2 AMI
      ___|\___|___|

https://aws.amazon.com/amazon-linux-2/
[ec2-user@ip-172-31-52-186 ~]$ sudo su -
Last login: Mon Feb 15 14:08:34 UTC 2021 on pts/0
[root@ip-172-31-52-186 ~]# kubuectl get nodes
-bash: kubuectl: command not found
[root@ip-172-31-52-186 ~]# kubectl get nodes
NAME                          STATUS   ROLES                  AGE   VERSION
ip-172-31-49-3.ec2.internal   Ready    <none>                 12d   v1.20.2
ip-172-31-52-186.ec2.internal   Ready    control-plane,master   12d   v1.20.2
[root@ip-172-31-52-186 ~]# kubeadm token list
[root@ip-172-31-52-186 ~]# kubeadm token create --print-join-command
kubeadm join 172.31.52.186:6443 --token zpp4vq.1utye931a9j01nyu     --discovery-token-ca-cert-hash sha256:c140d470d1bd13e48adf450bae48aae0d88a94f42472a5aa06842302d95f295e
[root@ip-172-31-52-186 ~]#
```

i-0484f53bd546fd94c (Kube master)

Public IPs: 54.236.51.47 Private IPs: 172.31.52.186

come to worker node again.

enter the above token directly in worker node.

```
[root@ip-172-31-93-26 ~]# kubeadm join 172.31.52.186:6443 --token zpp4vq.1utye931a9j01nyu     --discovery-token-ca-cert-hash sha256:c140d470d1bd13e48adf450bae
48aae0d88a94f42472a5aa06842302d95f295e
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

[root@ip-172-31-93-26 ~]#
```

verifying…

go to master

kubectl get node

```
[root@ip-172-31-52-186 ~]# kubectl get nodes
NAME                        STATUS   ROLES                  AGE    VERSION
ip-172-31-49-3.ec2.internal     Ready    <none>                 12d    v1.20.2
ip-172-31-52-186.ec2.internal   Ready    control-plane,master   12d    v1.20.2
ip-172-31-93-26.ec2.internal    Ready    <none>                 107s   v1.20.2
[root@ip-172-31-52-186 ~]#
```

Ip of node 2 is now mention as node in mater

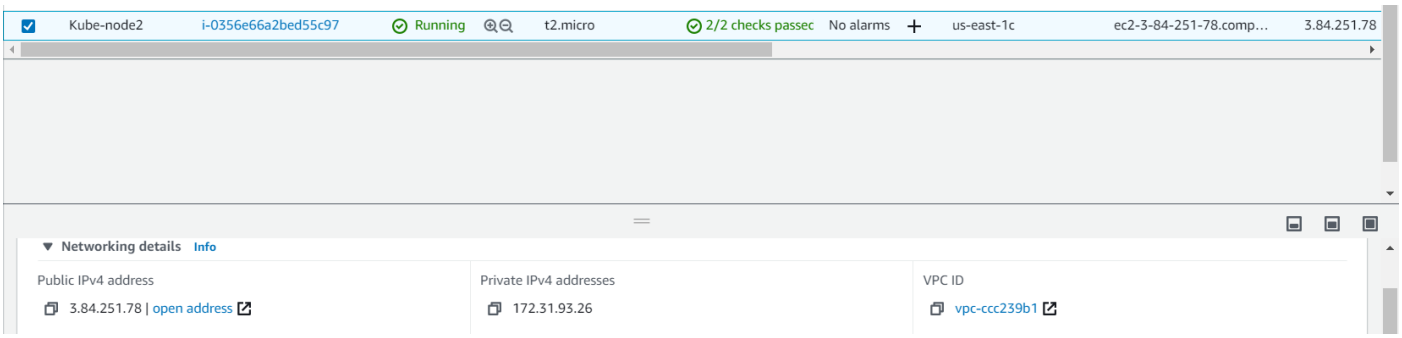

done……